

```
In [2]: import pandas as pd
import re
import numpy as np
import scipy
import itertools
import matplotlib
import matplotlib.pyplot as plt
from scipy.spatial.distance import cdist
from collections import Counter
from random import choice
```

```
In [3]: # Storing the training and test datasets into their respective dataframes
trained = pd.read_csv('train_clean.csv')
test = pd.read_csv('test_clean.csv')
```

```
In [4]: trained.head()
```

Out[4]:

	Unnamed: 0	Sentiment	Tweet
0	0	neutral	amsterdam ewr
1	1	negative	ITproblems link
2	2	positive	today staff MSP took customer service new le...
3	3	negative	yet receive assistance one agents securing ne...
4	4	negative	let change reservation online Im wasting time

```
In [5]: test.head()
```

Out[5]:

	Unnamed: 0	Sentiment	Tweet
0	0	neutral	jump DallasAustin market News
1	1	positive	Chicago seen seat A AA So far great ride On ...
2	2	negative	need bag bouncer Get together
3	3	negative	Hey Jetblue stranded entire plane supposed go...
4	4	negative	Big fail curbside baggage Pittsburgh charge ...

```
In [8]: #Training Data
train_unique = (list(set(trained['Tweet'].str.findall("\w+").sum()))) # Finding all unique words in training data
train_unique_words = len(train_unique)

#Test Data
test_unique = (list(set(test['Tweet'].str.findall("\w+").sum()))) # Finding all unique words in test data
test_unique_words = len(test_unique)

print("Unique words in Training Data: {}".format(train_unique_words))
print("Unique words in Test Data: {}".format(test_unique_words))
```

Unique words in Training Data: 12416
Unique words in Test Data: 5814

Feature Extraction

```
In [9]: #Training Data
train_matrix = [] # Forming a 2D matrix to store all training feature vectors

#Test Data
test_matrix = [] # Forming a 2D matrix to store all test feature vectors
```

```
In [10]: #Training Data: Extracting features and storing them into the training feature matrix
for sentence in trained['Tweet']:
    train_featurevec = []
    word = sentence.split()
    for w in train_unique:
        train_featurevec.append(word.count(w))
    train_matrix.append(train_featurevec)

#Test Data: Extracting features and storing them into the test feature matrix
for sentence in test['Tweet']:
    test_featurevec = []
    word = sentence.split()
    for w in train_unique:
        test_featurevec.append(word.count(w))
    test_matrix.append(test_featurevec)
```

```
In [11]: print("Shape of Training Matrix: ({0} , {1})".format(len(train_matrix),len(train_matrix[0])))
print("Shape of Test Matrix: ({0} , {1})".format(len(test_matrix),len(test_matrix[0])))
```

Shape of Training Matrix: (11680 , 12416)
Shape of Test Matrix: (2921 , 12416)

```
In [12]: #Calculating distances between every test instance with all the train instances.
dists = cdist(test_matrix,train_matrix,'euclidean')
```

```
In [13]: #Making an empty column in our test data for predicted labels.
test['Predicted Label'] = ''
dists.shape
```

```
Out[13]: (2921, 11680)
```

```
In [14]: #Function that takes a list and returns the mode of the list. If there are more t
def get_mode(l):
    counting = Counter(l)
    max_count = max(counting.values())
    return choice([ks for ks in counting if counting[ks] == max_count])
```

K Nearest Neighbors & Performance Measures

```
In [15]: # Making a general structure of our confusion matrix
cmatrix = pd.DataFrame({'Gold Positive': '', 'Gold Neutral': '', 'Gold Negative': ''},
                        index = ['Predicted Positive', 'Predicted Neutral', 'Predicted Negative'])

# Lists that will later store respective values for plotting
accuracy_list = []
recall_list = []
precision_list = []
F1_list = []
```

```

In [16]: def cmatrix_measures(k,dists,test,cmatrix):

    row_count = 0
    first_max = 0
    second_max = 0
    check_tie = False

    for ls in dists:
        sorted_distances_indices = np.argsort(ls) #Getting a sorted list of indices of
        knn_indices = []
        knn_indices = list(itertools.islice(sorted_distances_indices,k)) #Extracting

        knn_labels = []
        for i in knn_indices:
            label = trained['Sentiment'][i] #Extracting the label of the instance by index
            knn_labels.append(label) #Appending the label to our labels list.

        max_class = get_mode(knn_labels)
        first_max = max_class
        second_max = max(knn_labels)
        if first_max == second_max:
            check_tie = True
        predicted_label = max_class
        test['Predicted Label'][row_count] = predicted_label

        row_count += 1

#Creating a frequency DataFrame that will store value counts for each tuple of
testfreqdf = test.groupby(["Sentiment", "Predicted Label"]).size().reset_index()
testfreqdf

#Extracting values from the Frequency DataFrame and assigning to specific cells
cmatrix['Gold Positive']['Predicted Positive'] = testfreqdf['Frequency'][8]
cmatrix['Gold Neutral']['Predicted Positive'] = testfreqdf['Frequency'][5]
cmatrix['Gold Negative']['Predicted Positive'] = testfreqdf['Frequency'][2]
cmatrix['Gold Positive']['Predicted Neutral'] = testfreqdf['Frequency'][7]
cmatrix['Gold Neutral']['Predicted Neutral'] = testfreqdf['Frequency'][4]
cmatrix['Gold Negative']['Predicted Neutral'] = testfreqdf['Frequency'][1]
cmatrix['Gold Positive']['Predicted Negative'] = testfreqdf['Frequency'][6]
cmatrix['Gold Neutral']['Predicted Negative'] = testfreqdf['Frequency'][3]
cmatrix['Gold Negative']['Predicted Negative'] = testfreqdf['Frequency'][0]

#Extracting all three True Positives from the matrix to measure accuracy.
TP = cmatrix['Gold Positive']['Predicted Positive']
TNT = cmatrix['Gold Neutral']['Predicted Neutral']
TN = cmatrix['Gold Negative']['Predicted Negative']
total = testfreqdf['Frequency'].sum()
accuracy = ((TP+TNT+TN)/total)*100
accuracy = round(accuracy,2)
accuracy_list.append(accuracy)

#Extracting all recalls from the matrix to measure macroaveraged recall.
recall_pos = cmatrix['Gold Positive']['Predicted Positive']/cmatrix['Gold Positive']['Predicted Positive']
recall_neut = cmatrix['Gold Neutral']['Predicted Neutral']/cmatrix['Gold Neutral']['Predicted Neutral']
recall_neg = cmatrix['Gold Negative']['Predicted Negative']/cmatrix['Gold Negative']['Predicted Negative']
macroaveraged_recall = ((recall_pos+recall_neut+recall_neg)/3)*100

```

```
macroaveraged_recall = round(macroaveraged_recall,2)
recall_list.append(macroaveraged_recall)

#Extracting all precisions from the matrix to measure macroaveraged precision.
precision_pos = cmatrix['Gold Positive']['Predicted Positive']/(cmatrix.iloc[0,
precision_neut = cmatrix['Gold Neutral']['Predicted Neutral']/(cmatrix.iloc[1,0
precision_neg = cmatrix['Gold Negative']['Predicted Negative']/(cmatrix.iloc[2,
macroaveraged_precision = ((precision_pos+precision_neut+precision_neg)/3)*100
macroaveraged_precision = round(macroaveraged_precision,2)
precision_list.append(macroaveraged_precision)

#Extracting all F1_scores from the matrix to measure macroaveraged F1_score.
F1_pos = (2*precision_pos*recall_pos)/(precision_pos+recall_pos)
F1_neut = (2*precision_neut*recall_neut)/(precision_neut+recall_neut)
F1_neg = (2*precision_neg*recall_neg)/(precision_neg+recall_neg)
F1_score = ((F1_pos + F1_neut + F1_neg)/3)*100
F1_score = round(F1_score,2)
F1_list.append(F1_score)

print("\n\nConfusion Matrix with k = {}: \n".format(k))
print(cmatrix)
print("\nAccuracy with k = {}: {1}%".format(k,accuracy))
print("Macroaveraged Precision with k = {}: {1}%".format(k,macroaveraged_precision))
print("Macroaveraged Recall with k = {}: {1}%".format(k,macroaveraged_recall))
print("Macroaveraged F1-score with k = {}: {1}%".format(k,F1_score))
```

In [17]: *#Calling the function for each individual k*

```
cmatrix_measures(1,dists,test,cmatrix)
cmatrix_measures(3,dists,test,cmatrix)
cmatrix_measures(5,dists,test,cmatrix)
cmatrix_measures(7,dists,test,cmatrix)
cmatrix_measures(10,dists,test,cmatrix)
```

C:\Users\aj240\AppData\Local\Temp\ipykernel_112\4185092973.py:24: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
test['Predicted Label'][row_count] = predicted_label
```

Confusion Matrix with k = 1:

	Gold Positive	Gold Neutral	Gold Negative
Predicted Positive	263	107	205
Predicted Neutral	132	343	758
Predicted Negative	77	165	871

Accuracy with k = 1: 50.56%

Macroaveraged Precision with k = 1: 50.6%

Macroaveraged Recall with k = 1: 52.99%

Macroaveraged F1-score with k = 1: 48.82%

C:\Users\aj240\AppData\Local\Temp\ipykernel_112\4185092973.py:24: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
test['Predicted Label'][row_count] = predicted_label
```

Confusion Matrix with k = 3:

	Gold Positive	Gold Neutral	Gold Negative
Predicted Positive	273	124	235
Predicted Neutral	135	363	783
Predicted Negative	64	128	816

Accuracy with k = 3: 49.71%

Macroaveraged Precision with k = 3: 50.83%

Macroaveraged Recall with k = 3: 53.79%

Macroaveraged F1-score with k = 3: 48.39%

C:\Users\aj240\AppData\Local\Temp\ipykernel_112\4185092973.py:24: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
test['Predicted Label'][row_count] = predicted_label
```

Confusion Matrix with k = 5:

	Gold Positive	Gold Neutral	Gold Negative
Predicted Positive	272	133	250
Predicted Neutral	138	376	814
Predicted Negative	62	106	770

Accuracy with k = 5: 48.55%

Macroaveraged Precision with k = 5: 50.64%

Macroaveraged Recall with k = 5: 53.58%

Macroaveraged F1-score with k = 5: 47.51%

C:\Users\aj240\AppData\Local\Temp\ipykernel_112\4185092973.py:24: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
test['Predicted Label'][row_count] = predicted_label
```

Confusion Matrix with k = 7:

	Gold Positive	Gold Neutral	Gold Negative
Predicted Positive	275	121	247
Predicted Neutral	153	397	871
Predicted Negative	44	97	716

Accuracy with k = 7: 47.52%

Macroaveraged Precision with k = 7: 51.42%

Macroaveraged Recall with k = 7: 53.95%

Macroaveraged F1-score with k = 7: 47.18%

C:\Users\aj240\AppData\Local\Temp\ipykernel_112\4185092973.py:24: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
test['Predicted Label'][row_count] = predicted_label
```

Confusion Matrix with k = 10:

	Gold Positive	Gold Neutral	Gold Negative
Predicted Positive	275	112	252
Predicted Neutral	158	424	893
Predicted Negative	39	79	689

Accuracy with k = 10: 47.52%

Macroaveraged Precision with k = 10: 52.39%

Macroaveraged Recall with k = 10: 54.92%

Macroaveraged F1-score with k = 10: 47.42%

In []: